

LINEAR REGRESSION MACHINE LEARNING MODEL

This programme is to visualise the USA_HOUSE_PRICING data set and run a linear regression model for predicting the house prices. Before we start Data visualisation is important.


```
In [2]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

%matplotlib inline
```

```
In [3]: df = pd.read_csv('USA_Housing.csv')
```

```
In [4]: sns.distplot(df['Price'])
```


```
Out[4]: <matplotlib.axes._subplots.AxesSubplot at 0x10f6ee860>
```



```
In [5]: sns.heatmap(df.corr(),annot = True
)
```

*#this corr map is used to find the correlation between the attributes and the pricing so as to
#get an idea about the relationship between them.*

```
Out[5]: <matplotlib.axes._subplots.AxesSubplot at 0x10f776080>
```



```
In [6]: df.columns
```

```
Out[6]: Index(['Avg. Area Income', 'Avg. Area House Age', 'Avg. Area Number of Rooms',  
            'Avg. Area Number of Bedrooms', 'Area Population', 'Price', 'Address'],  
            dtype='object')
```

```
In [7]: x = df[['Avg. Area Income','Avg. Area House Age','Avg. Area Number of Rooms',  
              'Avg. Area Number of Bedrooms','Area Population']]
#separate the other attributes from the predicting attribute
```

```
In [8]: y = df[['Price']]
#separte the predicting attribute into Y for model training
```

```
In [9]: from sklearn.model_selection import train_test_split

#import model selection train test split for splitting the data into test and train for
#model validation.
```

```
In [10]: x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.4, random_state=101)
```

```
In [11]: from sklearn.linear_model import LinearRegression
```

```
In [12]: lm = LinearRegression()
#Loading the model constructor
```

```
In [13]: lm.fit(x_train,y_train)
#training or fitting the train data into the model
```

```
Out[13]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=1, normalize=False)
```

```
In [14]: print(lm.intercept_)

[-2640159.79685191]
```

```
In [15]: lm.coef_
#examining the co-efficients of the fitted model.
```

```
Out[15]: array([[ 2.15282755e+01,  1.64883282e+05,  1.22368678e+05,
  2.23380186e+03,  1.51504200e+01]])
```

```
In [16]: x_train.columns
```

```
Out[16]: Index(['Avg. Area Income', 'Avg. Area House Age', 'Avg. Area Number of Rooms',
  'Avg. Area Number of Bedrooms', 'Area Population'],
  dtype='object')
```

```
In [17]: cdf = pd.DataFrame(data=lm.coef_.reshape(5,1),index=x_train.columns,columns=['Coeff'])
```

```
In [18]: cdf
```


```
Out[18]:
```

	Coeff
Avg. Area Income	21.528276
Avg. Area House Age	164883.282027
Avg. Area Number of Rooms	122368.678027
Avg. Area Number of Bedrooms	2233.801864
Area Population	15.150420

```
In [19]: predictions = lm.predict(x_test)
```


```
In [20]: plt.scatter(y_test,predictions)
#to visualise the predictions and the test Y !! almost it is forming a linear line with less deviation
```

```
Out[20]: <matplotlib.collections.PathCollection at 0x113d992b0>
```



```
In [21]: sns.distplot((y_test-predictions))
```

Out[21]: <matplotlib.axes._subplots.AxesSubplot at 0x113d49d68>



```
In [22]: from sklearn import metrics
```

```
In [ ]: np.sqrt(metrics.sqrt)
#printing the error values of the model
```

```
In [24]: print('MAE:', metrics.mean_absolute_error(y_test, predictions))
print('MSE:', metrics.mean_squared_error(y_test, predictions))
print('RMSE:', np.sqrt(metrics.mean_squared_error(y_test, predictions)))
```

```
MAE: 82288.2225191
MSE: 10460958907.2
RMSE: 102278.829223
```

```
In [ ]:
```